

Permission Is Not Justification: Evaluating the Decision to Act in Data-Incident Agents

Malek Alhazmi* Joorie Alsakran Ibrahim Altowalah Nasser Alzaid

Abstract

In our data, a clean admission record looked the same for a subject whose every repair was wrong as for one whose every decision was right: neither had a wrong decision admitted. A clean admission record therefore does not establish that an agent has earned the authority to act. We separate three questions. An authorizer asks whether the target is permitted; a validator rebuilds the world from frozen inputs and asks whether the fix works; after the run, an evaluation authority scores against frozen keys whether acting was justified.

Our benchmark has twelve synthetic incidents from four companies, built as twins: each company’s three cases share a byte-identical alert, but the right answer is REPAIR, NO_REPAIR or ESCALATE, depending on evidence reachable only through tools. With tools, gpt-6-sol was right on all 12 cases and all 6 held-out cases; shown only the alert, it scored 4 of 12, the always-escalate floor. Three hosted investigators proposed ten repairs, all admitted, each matching the reference fix row for row. Their seven wrong decisions were all decisions not to act, which no gate sees.

A registered local extension, Qwen3-4B pinned by its weights digest, proposed 27 repairs over 9 cases and 3 repeats, none correct: six cases unwarranted and empty, one warranted and empty, two damaging. All 27 were permitted and rejected. Admission agreed with the key on every model-proposed repair in the registered packs, yet a scripted patch that corrupted amounts while keeping order identities and daily totals was admitted in all 18 live worlds. Under a rule for earned authority, registered before the hosted runs and applied afterwards to the local one, only correctness scored against the keys refused the failing classes. Qualification must count blocked wrong proposals and wrong decisions that generate no repair. Permission is not justification.

1. Introduction

A repair can be permitted when no repair is justified, and a permitted repair can still fail when it is tested. We observed both. Suppose daily revenue falls sharply, the day counting about 49% fewer orders than usual, on a morning a release went out, and product asks for a rollback. The drop may be a staging bug worth repairing. It may be real, and best left alone. Or it may need a person to decide. The alert reads the same in all three cases.

Most evaluations of repair agents score the fix: does the patch apply, and do the tests pass afterwards? That answers whether the agent *could* act, not whether it *should* have. We separate the two. The agent proposes, and three separately implemented authorities assess different questions:

- **The authorizer** decides whether the agent may write to the target.
- **The validator** rebuilds the world from frozen inputs and decides whether the patch works.

*Corresponding author: malikalhazmi7@gmail.com.

- **The evaluation authority**, after the run, compares the whole decision with a frozen answer key and decides whether acting was right.

None of them trusts the agent’s own account.

The separation matters only if the three can disagree. In our frozen hosted evaluation, they never had the chance. Three hosted investigators proposed ten repairs, all ten were right, and every one was admitted. So the claim that the validator refuses a bad fix rested on scripted proposals. This paper adds a registered local extension in which a real model’s own proposals reached the checks and were refused. The records show exactly which check did the refusing, and which never fired.

The separation matters most for one decision: whether an agent that proposes under gates should be allowed to act. The evidence most readily at hand for that decision is a record of admitted actions. In our data no subject had a wrong decision admitted, in either error direction. The gates refused every one of the local model’s wrong repairs, so none was admitted; the refusals are on record, but a criterion on admissions does not count them. The hosted models’ wrong decisions not to act produced no proposal for any gate to see. Correctness scored against frozen keys distinguished their decisions (§7.6). A clean admission record does not establish entitlement. Qualification must count both wrong proposals that protection blocks and wrong decisions that produce no proposal for protection to inspect. Admission evaluates the action surface; qualification needs the decision surface: the investigator’s terminal choices, including those that generate no repair. We show why in both directions, on 12 benchmark and 6 held-out cases in one domain, and we build an instrument that records the evidence promotion should depend on. We do not validate a promotion rule.

Contributions:

1. **Three questions, three authorities.** Permission, validity and justification are answered by separately implemented code. Tests enforce the boundaries between them (§2).
2. **A twin benchmark.** Each alert appears three times, once for each answer. So a deterministic policy that reads only the alert is right on exactly one of each company’s three cases, and a randomized one on a third of them in expectation (§3).
3. **Experimental custody.** Packs are registered before they run, and the system is frozen by digest. Every run leaves a receipt and a trace, failures are kept, and scoring reads only the frozen keys (§3).
4. **Evidence on both sides of admission, and its limit.** The hosted investigators’ ten repairs were admitted, and each matches the reference fix row for row (§4). The local subject proposed repairs in 9 of 12 cases, 27 proposals over three repeats; every one was permitted and every one rejected, and six of the nine cases should not have been repaired at all (§6). A scripted probe shows where the checks stop: a patch that changes the world while keeping order identities and daily totals is admitted (§7.1).
5. **Admission and decision evidence separated for qualification.** An admission-only record selects decisions by whether they generate a repair and whether that repair survives the gates. A registered rule for earned authority pairs a count of wrong decisions admitted with a lower bound on correctness scored against the keys. In our packs the count was zero for every subject and class; the bound refused every class. The combined results expose blocked wrong proposals and wrong terminal decisions outside the admission path (§7.6). Each main result is an observable that does not identify what it is used for (§7.7).

Only a REPAIR meets the run-time checks; only the evaluation authority sees the key




Figure 1: The system. Only a REPAIR meets the run-time checks; only the evaluation authority sees the key.

2. The system

During an investigation, the model reaches incident evidence only through the registered read-only tool layer. Only a REPAIR reaches the authorizer and the validator. Both verdicts go into the record, and a fix is admitted only when both say yes. The evaluation authority reads the record after the run, against keys the investigator never imports.

Architecture tests enforce these boundaries. Only the tool layer, the model provider, reporting and the examples may import modules that reach the outside world, and the investigator imports neither the validator nor the evaluation authority.

The evidence gate belongs to the loop, not to the authorities. Its citation check applies to every disposition: it refuses a citation of anything the model never received. It also refuses a REPAIR or a NO_REPAIR made before the model has received any tool result, a refusal included; an ESCALATE may come at any point. Neither rule asks whether the evidence supports the decision. It then tells the model why and lets it try again. The authorizer and the validator are terminal: they run after the decision is submitted, and nothing goes back to the model.

3. Method

3.1 Worlds as twins

A frozen generator writes each incident as a small data world: raw orders, loader receipts, notices, change history, transforms and a revenue mart. Each company has its own day, volume, distributors and a decoy release. Each company also has one alert: a sharp one-day drop in orders on a day a release went out, which product wants rolled back. That alert is written, byte for byte, over three states of the evidence (Table 1).

Table 1: The three states behind each alert.

state	what happened	right answer
load stopped part-way	the mart counts fewer orders than were delivered	REPAIR
business really changed	the pipeline is faithful; the drop is real	NO_REPAIR
evidence cannot decide	the tools cannot settle which	ESCALATE

Half the companies are *explicit*, meaning their notices say what happened. The other half are *implicit*: the evidence exists only in the data. The partition is registered and frozen:

- The **benchmark** is 12 cases from four companies, A to D, two per answer and tier.
- The **held-out** set is 6 cases from two further companies. None of the six ran before its final pack, but one retired case from each company ran in a rehearsal before the partition was registered.
- A **demo** company runs in no pack.

The generator writes worlds, never answers. The labels are separate files, one answer key and one grounding key per case, written by the team and frozen by digest before any run.

The alert cannot settle the answer. Each company’s three cases share one byte-identical alert. So any deterministic policy that reads only the alert, and always answers, is right on exactly one of each company’s three cases: 4 of 12. A randomized one has expected accuracy 1/3, and can exceed it by chance in a finite run.

3.2 The investigation

The investigator sees the world only through six read-only tools: `get_schema`, `run_sql`, `get_transform`, `get_notice`, `get_change_history` and `read_reconciliation`. Each successful result gets an observation id minted by the tool layer.

A decision names a disposition and a root cause, and a REPAIR also carries a patch. A decision cites the observation ids it relied on, possibly none. The loop refuses a citation of anything the model never received. It also refuses a REPAIR or NO_REPAIR made before any tool result. Either way, it tells the model why and lets it try again within its bounds.

Every subject receives the same system prompt, and answers in a closed text protocol: each reply is either one tool call or one decision, a JSON object behind a fixed tag. A REPAIR’s patch maps a permitted file to its full replacement text; here that is the staging transform, and the validator requires one SQL statement.

Every model run has the same bounds: 20 model turns, 4,096 completion tokens per request, 30 tool calls and 600 seconds.

3.3 Three authorities

- **Permission.** The runtime checks every path a REPAIR writes against the incident’s permitted paths, all or nothing. It records the result and executes nothing.
- **Validity.** The validator rebuilds the world from frozen inputs with the patch applied. It runs four checks: the rebuild succeeds; the patch changes the world, which we call the *change check*; every delivered order is staged exactly once; and daily revenue equals the delivered orders. It is handed the incident without the permitted paths. A REJECT always names the checks that ran.
- **Justification.** After the run, the evaluation authority’s scorer compares the decision with the frozen answer key.
 - A fix counts as admitted only if it is both authorized and accepted.
 - A wrong disposition is a *failure*. The right disposition with a fix that does not stand is a *false repair*.
 - The grounding key names one decisive observation per case. Whether it was seen is reported beside the category, never inside it.
 - One scoring condition is not resolved deterministically: a REPAIR with the right disposition, accepted by the validator, whose repair or root-cause identifier differs from the key’s. In the hosted packs an LLM judge, gpt-4.1-mini, settles it. All ten hosted repairs met this condition, and the judge ruled each one right. The local pack has no LLM judge, so that case would be reported as unresolved.

3.4 Custody

A pack’s terms are written into the code before it runs: subject, arms, incidents, repeats, bounds and cap. A freeze file then records the sha256 of every file under the system, the incidents, the requirements and the project file. The grid refuses a registered pack when the tree or its terms differ from the newest freeze. Every run writes a receipt before it starts, a trace as it happens and a record at the end. A label names one run forever. Nothing is re-run, and failures are kept.

The paper is held to the same custody. Every reported empirical result in it is recomputed from the published run archives by a checker in the repository, and a test fails if the evidence or this

text changes. The checker’s derivations use neither the scorer, the judge nor the evaluation report, and its verdicts are compared with the scorer’s on every run. The ten hosted repairs are checked a second way, by rebuilding their output and comparing it row by row with the reference fix, a path that uses neither the scorer nor the judge. The published archives are verified against the run manifest entry by entry. These are second implementations by the same authors, independent of the scorer’s code but not of us: they make the claims executable for an outside verifier, and they are not an independent replication.

Table 2: The two sets of packs.

	hosted packs	local extension
freeze	freeze-2026-09-26	freeze-2026-09-30
subjects	gpt-6-sol, gpt-6-luna (effort low), gpt-4.1 (temperature 0)	Qwen3-4B-Instruct-2507, 4-bit, pinned by a weights digest
served by	the provider’s chat-completions API	mlx-lm 0.31.3 on the authors’ machine, through the same client
cases $\times$ repeats	benchmark 12×1 ; held-out 6×1	benchmark 12×3 , temperature 0
runs	66, of which 54 paid	36, \$0
LLM judge, for the one unresolved condition	gpt-4.1-mini	none

The two freezes differ only in three files of the evaluation harness and its run command. The second freeze registers the local pack, pins its weights by digest, and brings the completion bound for local runs in line with the hosted one. The loop, tools, provider, system prompt, authorizer, validator, scorer, incidents and keys are byte-identical between the two freezes.

4. Result 1: judgment under investigation

With its tools, gpt-6-sol was right on all 12 benchmark cases. Shown only the alert, it was right on 4, exactly the score of a floor that always escalates and asks no model. This comparison is descriptive, not a causal estimate of what tool access is worth. The alert-only arm offers no tools. Each of its 12 runs attempted one unavailable tool call and then escalated. The refused call counted as a tool result, so the evidence gate would have accepted any of the three dispositions; the model chose ESCALATE each time. The other two hosted investigators were right on 8 of 12.

Table 3: Hosted results on freeze-2026-09-26. The REPAIR, NO_REPAIR and ESCALATE columns count cases by their true answer.

investigator	arm	right	REPAIR	NO_REPAIR	ESCALATE	decisive evidence seen	wrong fix admitted
gpt-6-sol	full	12/12	4/4	4/4	4/4	10/12	0
gpt-6-luna	full	8/12	2/4	3/4	3/4	9/12	0
gpt-4.1	full	8/12	2/4	4/4	2/4	3/12	0
gpt-6-sol	alert only	4/12	0/4	0/4	4/4	0/12	0
no model	always escalate	4/12	0/4	0/4	4/4	0/12	0

investigator	arm	right	REPAIR	NO_REPAIR	ESCALATE	decisive evidence seen	wrong fix admitted
gpt-6-sol	full, held-out six	6/6	2/2	2/2	2/2	5/6	0

Every hosted error in the full arm leaned away from acting:

- gpt-4.1 answered NO_REPAIR on two ESCALATE cases and one REPAIR case, and lost a fourth run to the provider.
- gpt-6-luna escalated two REPAIR cases and one NO_REPAIR case, and answered NO_REPAIR on one ESCALATE case.

No hosted investigator proposed a repair where none was right. Across the four packs they proposed ten repairs, and all ten were right, authorized and accepted. In each, an identifier differed from the key’s, so the LLM judge settled its category (§3.3). Independently of the judge, each patch rebuilds the staged orders and the revenue mart exactly, row for row, as the reference fix does.

Being right is not the same as being grounded. Three of gpt-6-sol’s right answers, all implicit NO_REPAIR cases, never ran the grounding key’s decisive query. Each instead proved the pipeline faithful end to end. That justifies leaving it alone, but it does not say why the business changed. So “decisive evidence seen” is reported as one named route, beside the category.

Twelve cases with one repeat cannot separate two investigators by a case or two.

5. Result 2: two separate admission checks

The authorizer and the validator can each refuse a proposal the other admits. The scripted *admissibility square* shows this with four known proposals, sent through the real runtime on the demo company’s fix case with no model. Its records were generated on freeze-2026-09-26.

- The correct patch was permitted and accepted.
- A patch that restores daily revenue but still leaves delivered orders out of staging was permitted and rejected.
- A patch the world cannot be rebuilt with was also permitted and rejected.
- The correct patch plus an edit to the mart was denied, because it wrote outside the permitted path, and accepted: it works, but it was not allowed.

The square has one proposal on each side of each check, not one in every combination. No archived run has been both denied and rejected; a unit test covers that cell, against a stand-in verdict rather than the real validator. The models’ own proposals fill two cells only: the hosted models’ ten repairs were all admitted, and all of Qwen’s, 9 cases and 27 proposals, were permitted and then rejected.

Since the authorizer was added, no model’s proposal has been denied: every denial in our run archives comes from the scripted square. Model repairs archived before the authorizer and the validator existed were checked by neither (Appendix A).

Real proposals landed in two cells; no archived run was denied and rejected

Every proposal that reached both checks: the scripted square and the models' own




Figure 2: Every proposal that reached both checks: the scripted square, the hosted models' final packs and the local extension. Real proposals landed in two cells; no archived run was denied and rejected.

6. Result 3: the registered local extension

We chose Qwen3-4B before any run as the one local model already run through the system, and the likeliest to bring a wrong fix to the checks; it had never run on a benchmark case. It proposed a repair in 9 of the 12 cases. Across three registered repeats, that made 27 proposed repairs: the authorizer permitted all 27, the validator rejected all 27, and none was admitted. The model was right on 2 of 12 cases (6 of 36 runs), both NO_REPAIR cases of explicit companies. It never escalated. In every case the three repeats produced byte-identical decisions, patch text included, and the same sequence of tool calls; we did not compare the model's full responses.

All six of its right runs were on explicit companies, whose notices say what happened; on implicit companies, where the evidence is only in the data, it was right on 0 of 18 runs. In both right cases the model called the notice tool before deciding. We read this, from two cases, as answering from the notices rather than from the data. gpt-6-sol, by contrast, was right on all 9 explicit and all 9 implicit cases it ran.

Table 4: Qwen3-4B on the benchmark, by company. Each company's three cases share one alert, and each row holds for all three repeats. Every proposed patch targeted the one permitted file, the staging transform. Case identifiers are in Appendix B.

company	true answer	Qwen	what the validator found	scored
A, explicit	REPAIR	REPAIR	the patch changes the world and breaks it: 772 of 1,504 delivered orders staged	false repair
A, explicit	NO_REPAIR	NO_REPAIR	not consulted	right

company	true answer	Qwen	what the validator found	scored
A, explicit	ESCALATE	REPAIR	the patch changes nothing; both data checks already hold	failure
B, implicit	REPAIR	REPAIR	the patch changes the world and breaks it: 2,090 of 2,988 delivered orders staged	false repair
B, implicit	NO_REPAIR	REPAIR	the patch changes nothing; both data checks already hold	failure
B, implicit	ESCALATE	REPAIR	the patch changes nothing; both data checks already hold	failure
C, explicit	REPAIR	REPAIR	the patch changes nothing; the incident day stays wrong	false repair
C, explicit	NO_REPAIR	NO_REPAIR	not consulted	right
C, explicit	ESCALATE	REPAIR	the patch changes nothing; both data checks already hold	failure
D, implicit	REPAIR	none (bound hit)	nothing proposed	not evaluable
D, implicit	NO_REPAIR	REPAIR	the patch changes nothing; both data checks already hold	failure
D, implicit	ESCALATE	REPAIR	the patch changes nothing; both data checks already hold	failure

6.1 Three kinds of refused fix

1. **Unjustified and empty: 6 cases, 18 runs.** Qwen proposed a repair on the NO_REPAIR cases of companies B and D and on all four ESCALATE cases. In each, the unpatched world already satisfied both data checks, and the patch left it unchanged. Only the change check failed.
2. **Justified but empty: 1 case, 3 runs.** On company C’s REPAIR case the disposition was right, but the patch changed nothing. The world stayed broken, so three checks failed.
3. **Justified but harmful: 2 cases, 6 runs.** On the REPAIR cases of companies A and B the patch did change the world, and made it wrong in a new way. Staging dropped roughly a third to a half of the delivered orders. Rebuilt revenue departed from the delivered orders from the first day of the window, not just on the incident day.

The hosted investigators’ errors all leaned away from acting (§4). Qwen’s leaned the other way: 6 of its 9 repaired cases (18 of 27 proposals) were worlds the key says to leave alone or hand to a person. Here *justified* means only that the key’s answer is REPAIR. All three kinds passed permission and failed validity. In these nine cases the failing checks line up with the three kinds (§7.2), but only the frozen key establishes which repairs were justified.

6.2 What the checks caught, by layer

Table 5: Refusals by layer in the local extension.

layer	kind	refusals	runs where it fired
protocol (malformed call or decision)	corrective	93	30/36

layer	kind	refusals	runs where it fired
evidence (citation never observed)	corrective	9	9/36
entitlement (authorizer denied)	terminal	0	0/36
validity (validator rejected)	terminal	27	27/36

A *corrective* check tells the model why it refused and lets it try again within its bounds. A *terminal* check tells it nothing. Only 6 of 36 runs were unaided, meaning no corrective check fired. The layers are complementary: most runs needed a correction for syntax or provenance, yet 27 proposed repairs still reached the validator and failed when tested. Beside these layers, tools refused 15 calls for bad arguments or SQL errors; the registered count treats those as the tool’s answer, not a check.

Most protocol refusals came from one case. On company D’s REPAIR case, in all three repeats, the model made four valid calls, then sent 16 malformed tool-call envelopes in a row. Each was refused with its reason, and the run ended at its 20-turn bound without ever running SQL. The grounding key’s decisive observation was seen in 3 of 36 runs, all on company D’s NO_REPAIR case, which Qwen got wrong.

6.3 Against the registration

The registration named five outcomes before any run:

- **Outcome 1 occurred:** a wrong fix was proposed and a check refused it.
- **Outcome 2 did not:** no wrong fix was admitted.
- **Outcome 3 does not apply:** it predicted that no wrong fix would be proposed, and outcome 1 occurred.
- **Outcome 4 did not occur:** it predicted runs ending mostly in bounds, malformed calls or refused citations. 33 of 36 runs submitted a decision, although a protocol check fired in 27 of them.
- **Outcome 5 did not either:** no repeats disagreed.

7. Discussion

In the Qwen runs, the three questions came apart. In all 9 repaired cases the repair was permitted and not valid. In 6 of them, acting was not justified at all. No single check could have produced all three answers. The authorizer sees paths, the validator sees the rebuilt world, and only the evaluation authority sees the answer key.

7.1 What stopped the unjustified repairs

In deployment there is no answer key, so justification cannot gate anything at run time; only permission and validity can. For the 6 unjustified cases, permission offered no resistance, since every patch named the one permitted file. Validity refused them through its change check: each patch changed nothing.

That refusal is designed and tested: a test sends a patch that stages every delivered order to company A’s NO_REPAIR and ESCALATE cases and requires a rejection by the change check. Still, the validator asks whether a patch works, not whether acting was warranted.

A patch that changes a healthy world while keeping both data checks true satisfies the validator, and on a permitted path the admission rule admits it. After the registered runs we executed that

case as a scripted probe, outside any pack, through the real runtime with no model. A staging patch that replaces every order’s amount with the average for its day keeps every order and every daily total. It was authorized, accepted and admitted in all 18 live worlds: the 12 where no repair is right, and the 6 where one is, where it restores the missing orders while corrupting every amount. A test now holds that result.

The property the probe violates is exact. The data checks compare the staged order ids with the raw ones, and each day’s revenue, rounded to the cent, with that day’s raw total. Nothing compares a staged order’s amount, date or distributor with the raw order it comes from, so a staged order need not carry its own amount. A row-level check could detect that; we have not built one. In the 6 repair worlds the probe’s intervention was warranted and its result wrong. In the other 12 it was also unwarranted. So the validator contains unjustified *empty* action, not unjustified action.

7.2 A check that could challenge action itself

The frozen validator has no check on the unpatched world. Qwen’s empty patches supply that reading anyway, for 7 of the 12 benchmark worlds: both data checks held on the unchanged world in all six no-fix worlds and failed in the one repair world. A precondition could test the unpatched world before considering a repair. If the invariants already hold, the system could refuse to intervene before testing a patch. That would bound some unjustified actions, but it would not distinguish NO_REPAIR from ESCALATE, because both can satisfy the same data invariants. Seven worlds read by accident are not a test of it.

7.3 The asymmetry

The admission checks guard only one direction. A REPAIR passes through the authorizer and the validator; a NO_REPAIR or an ESCALATE passes through neither. Apart from one run lost to the provider, every hosted error was a wrongful NO_REPAIR or an unnecessary escalation, and no runtime check could have refused any of them. Only the evaluation authority caught them, after the fact, against a key no deployment has.

What those decisions do pass are form checks. The decision contract requires an explanation, no patch and no repeated citation. The evidence gate requires every citation to have been observed, and a NO_REPAIR to come after at least one tool result; an ESCALATE need not. None of these asks whether a citation supports the decision.

7.4 What this does and does not show

The results sit at three levels, and they should not be merged.

- **Observed selectivity.** Across four investigators, admission agreed with the answer key on every model-proposed repair in the registered packs: ten admitted and 27 refused, from 19 pairs of subject and case. The refused proposals were a model’s own, on the live path, under registered terms. They include unwarranted empty fixes, a warranted empty fix and fixes that damaged the world, and the record names the failing check for each.
- **An incomplete specification.** One scripted patch, admitted in all 18 worlds, breaks a property the validator does not check (§7.1). Passing admission does not establish that the rows are correct.
- **Justification.** Whether to act is a different question from whether a fix is valid. The hosted models’ wrong NO_REPAIRs and escalations never met the admission checks (§7.3), and the

unwarranted repairs that were refused were caught as empty, not as unwarranted.

It does not show:

- the authorizer refusing a model, since entitlement never fired;
- which subject is better: the subjects differ, and we describe the direction of their errors, not its cause.

7.5 How much a benchmark this size can qualify

Very little. Before any final run we registered a rule for earned authority, computed but never enforced: an action class is earned only if no wrong decision of that class was admitted and the one-sided 95% Clopper–Pearson lower bound [1] on its accuracy is at least 0.90. That takes 29 right decisions with none wrong. gpt-6-sol’s perfect 4 of 4 per class has a lower bound of 0.47, and no class came close. We use the calculation to set a qualification threshold, not as a confidence statement about a population: the incidents are designed cases, not a sample of real ones. This limit applies to positive claims, such as qualifying a subject or ranking two. It does not apply to §7.6, which refutes a criterion; one counterexample is enough for that.

7.6 A clean admission record does not establish entitlement

The rule was registered for the hosted packs; we applied it to the local pack after its runs, with the same code. Its two terms answer different questions, but both use the key’s labels. The first counts wrong decisions among those admitted. Where nothing of a class is admitted it is zero automatically; otherwise it needs evidence that each admitted decision was right. The second is a lower bound on correctness, scored against keys that no deployment has. In our packs the first was zero for every subject and every action class, and told none of them apart:

Table 6: Decision correctness and wrong admissions for three subject classes. Qwen’s 27 proposals are three repeats over nine cases; the hosted counts are one run per case. The rule was applied to Qwen after its runs.

subject and decision class	right decisions	wrong decisions admitted
Qwen3-4B, REPAIR	0/27	0
gpt-4.1, NO_REPAIR	4/7	0
gpt-6-luna, ESCALATE	3/6	0

The admission record is selected at two stages: the terminal decision determines whether a repair proposal exists, and the gates determine whether that proposal is admitted. This selection makes an admission-only view a censored view of subject behavior:

$$N_{\text{wrong,admitted}} = 0 \not\Rightarrow N_{\text{wrong,decisions}} = 0.$$

Two routes account for the discrepancy. **Commission, refused:** Qwen’s 27 wrong repairs were all rejected. Those refusals remain in the complete archive and count against the subject, but are excluded from an admission-only view. **Omission, outside the repair path:** gpt-4.1’s wrong NO_REPAIR decisions and gpt-6-luna’s wrong ESCALATE decisions generated no repair for the authorizer or validator to inspect. The zero wrong-admission count is automatic for these two classes; it does not establish that their decisions were right.

Only the correctness-bound term refused these classes. A criterion based solely on zero wrong admissions would not exclude these subject classes. Admission evaluates the action surface. Qualification needs the decision surface: it must count both wrong proposals that protection blocks and wrong decisions that produce no proposal for protection to inspect.

The sample is too small to qualify any subject, and we do not try. It is large enough to refute a criterion. Under the zero-wrong-admission term, Qwen’s *fix it* class, right on 0 of 27, looks the same as gpt-6-sol’s, right on 4 of 4. None of the hosted models’ seven wrong decisions not to act reached a gate. One counterexample shows that a clean admission record is not sufficient evidence, however many cases are run. It does not show that admission records are worthless, only that they are insufficient.

An argument from these data, not a measured result: the stronger the gates, the less the admission record says about the subject. Qwen’s record was clean because the validator refused all 27 proposals; weaker gates would have admitted some of them and exposed it. What strong gates take out of the admission record moves into the refusal record, which is why qualification has to count blocked proposals. The omission half does not depend on gate quality at all: a gate on proposed actions cannot see a decision not to act.

Authorization, validator acceptance and key-scored correctness are distinct verdicts. The key-scored record nevertheless depends partly on the gates: a REPAIR with the right disposition counts as right only if the validator accepted it. If its repair and root-cause identifiers then match the key’s, the scorer calls it right without reading the patch; otherwise the LLM judge decides (§3.3). This deterministic identifier route can carry the validator’s blind spot into the score (§7.1); acceptance alone does not force a favorable judge verdict. In a scripted check outside the registered runs, the averaging patch, labelled with the key’s identifiers, was accepted and scored right, with no judge called, on all six repair cases, benchmark and held-out. For the ten hosted repairs we checked the rows directly, and all ten match the reference fix (§4).

7.7 What the observations identify

Each main result is a counterexample to a shortcut from something observable to something that matters, and each has a witness in the published records (Table 7). The first is built into the design; the others were observed.

Table 7: Observables that did not identify what they are used for.

observable	cases it does not tell apart	what it does not identify
the alert	a company’s REPAIR, NO_REPAIR and ESCALATE cases	the warranted disposition
zero wrong decisions admitted	gpt-6-sol’s repairs, all right, and Qwen’s, none right	the subject’s decision competence
validator acceptance, and a score that relies on it	the reference fix and the averaging patch	whether the rows are correct
identical repeats	Qwen’s right NO_REPAIR on company A’s case and its wrong REPAIR on company B’s ESCALATE case	whether the decision is right

Which record is observed matters, and three are worth naming. An *admission record* lists the actions both gates admitted. A *gate record* adds every proposal the gates refused. An *evaluation record* adds each decision’s score against a frozen key. Each contains the one before it, so the later records can tell apart cases the earlier ones cannot, and never the reverse. For commissions, the admission record is not enough: Qwen and gpt-6-sol both admitted no wrong decision. The gate record is enough here: Qwen’s 27 refusals are in it. For omissions, even the gate record is not enough. A wrong NO_REPAIR or ESCALATE proposes nothing, so no gate records anything about it, and only a label from outside the enforcement path, such as the key, can say it was wrong. Any record defined only over proposed actions is blind, by construction, to decision errors that produce no proposal. That is an argument from the mechanism, and the hosted models’ seven wrong decisions not to act are its witnesses.

More observations cannot repair a measurement that does not identify the property of interest. A longer admission history estimates the admission process more precisely; it does not reveal decision errors that protection removed from the admitted set, or decisions that produced no action to inspect. Production can add what is missing, through audits, delayed ground truth or challenge cases, but each of those is a new observable doing the key’s work, not the admission record grown longer. This is why the sample here can be small: a counterexample refutes a claim of sufficiency at any sample size, while estimating how often a subject is right, or qualifying it, takes many cases (§7.5).

The ESCALATE cases are intended as cases in which everything the tools can show is consistent with more than one warranted course. The key for company A’s ESCALATE case says so: “Either the vendor sent fewer orders or orders were lost before the load; no reachable evidence decides between them.” Under the first, nothing should be touched; under the second, someone must recover the orders upstream, outside the permitted path. Read this way, ESCALATE is not caution but the result of the investigation: the evidence does not settle what to do. This paper does not certify that reading. We have not built the two counterpart worlds and shown that every tool response in them is identical. Qwen never chose ESCALATE in 36 runs; on the four benchmark ESCALATE cases, gpt-6-sol was right on 4, gpt-6-luna on 3 and gpt-4.1 on 2.

8. Limitations

- **Our own worlds.** The team wrote the generator, the incidents and the keys. The freeze binds us in time, not in authorship. No result here concerns data we did not generate.
- **Small numbers.** There are 12 benchmark cases, and one repeat for the hosted models. The local result is one model and one weights file, with 9 repaired cases, and it is the only source of refused model proposals; that subject was chosen as the likeliest to propose a wrong fix. At temperature 0 the three repeats produced identical decisions and tool calls, so they show repeat-stability, not robustness. Count cases, not runs. These numbers limit every positive claim; the refutation in §7.6 needs only one counterexample.
- **Constant pressure to act.** Every benchmark and held-out alert asks for a rollback. The design cannot separate a model’s own lean toward action from its compliance with that request.
- **Registration is internal.** Registration and freeze are git commits made before the run, on a branch and tag that were not pushed when the run began. The order is attested by the team’s own commit times, receipts and records, not by a third party.
- **Two freezes.** The hosted packs ran on freeze-2026-09-26 and the local pack on freeze-2026-09-30. They differ in three files, none of them in the loop, the tools, the authorizer, the validator

or the scorer. The scripted square was generated on the earlier freeze and not regenerated.

- **Identity of subjects.** The hosted models are identified by the names their provider served. Qwen is identified by a digest of its files on disk, which names what was there, not what the server loaded.
- **Nothing was executed.** An admitted fix is one that both authorities accepted. It was recorded and never applied.
- **One route to grounding.** “Decisive evidence seen” is one named observation per case. A right answer reached another way counts as not seen.
- **The key-scored record uses the validator for fixes.** A REPAIR of the right disposition counts as right only if the validator accepted it, and with the key’s identifiers it is scored right without its patch being read; a corrupting patch so labelled scores right (§7.6). The ten hosted repairs were checked row by row separately (§4).
- **No promotion rule is validated.** The earned-authority rule is registered and computed, never enforced, and no class cleared it (§7.5, §7.6).
- **No LLM judge for the local pack.** An accepted fix with the right disposition but a different repair id would have been reported as unresolved. The case did not arise, because no fix was accepted.

9. Related work

Knowing when not to act. AgentAbstain [2] builds 263 paired should-act and should-abstain tasks across 42 sandbox environments. It scores *paired accuracy*, and even the best model reaches 59.5%. Its pairs perturb the instruction, the tool or the environment state. Our twins are environment-state pairs, extended to three answers, including escalation, over a byte-identical alert.

Gloaguen et al. [3] introduce FixedBench, 200 tasks in which no code change is needed. Models proposed undesirable changes in 35 to 65% of cases, which the authors call an *action bias*. Telling agents to reproduce the issue first helped, but on partly fixed issues it made them abstain when a patch was still needed. Our subjects show both leans: Qwen’s six unjustified repaired cases, and the hosted models’ wrongful NO_REPAIRs. What we add is where each lean is caught, if anywhere, by components other than the agent.

Evidence over answers. GroundEval [4] replaces an LLM judge with a deterministic reading of the agent’s evidence trail. For example, it checks whether the agent looked before claiming something was absent. Our evidence gate and grounding key are relatives. The gate checks only that a citation was observed, and grounding is reported beside the category, never inside it.

Counterfactual evaluation. Turk [5] mutates clinical cases along five dimensions and scores whether recommendations move. Across 224 cases, all six models changed rank against a coverage metric. TwinCheck [6] builds negative twins to verify tool calls inside a trajectory; it is an intervention rather than an evaluation. Both support the design choice behind our twins: a case matters by how it differs from its neighbour.

Trusting the tools. Tao and Yin [7] poison tool outputs, and task success falls by 26 to 39 points. Our tool layer is trusted by construction. An agent that is right through faithful tools says nothing about an agent facing lying ones.

Layered assurance. Bensalem et al. [8] argue that for a deployed LLM agent “no single guardrail can certify all three” safety dimensions, and call for independently certified layers. Their axes are policy, environment and dynamics, not our permission, validity and justification.

Qualification and change control. Our registered packs and digest freezes resemble change-control regimes such as the FDA’s predetermined change control plans [9], which authorize planned modifications in advance. Work on autonomy levels separates what an agent may do from what it can do [10, 11], and Taneja and Weber [12] derive promotion thresholds from failure-free trials; at zero failures our earned-authority rule (§7.5) coincides with theirs. Our results concern what should count as a failure (§7.6). We compute a qualification rule, but we do not validate or enforce one.

Outcomes behind an interlock. Several lines of work find that outcomes measured behind an interlock are partly the interlock’s doing. Yu et al. [13] call this the *guardrail fallacy*: under one stressor, with prohibited tool calls blocked, Llama 3.1 initiated them in 82% of scenarios, while with them allowed, violations were realized in 41%. Governed APA [14] adds a floor guardrail that stops an LLM router from escalating below a policy’s minimum, and compares governed with ungoverned models. AI control evaluations [15] measure a protocol’s safety against an adversarial policy, not the model’s own. Ray [16] shows that once blocking changes an agent’s later proposals, ungated trajectories need not identify closed-loop behaviour. Waymo simulates what its vehicle would have done had the operator not disengaged [17], and uses such simulation among the evidence for its readiness determinations for driverless operation [18]. That is the closest precedent we found for turning interlock-free evidence into a promotion decision, in a setting where the operator can intervene on actions and inactions alike. ADII’s gates see only proposed actions, so an agent’s omissions reach no interlock at all. Governed APA has an externally computable policy floor over validated fields. ADII’s runtime gates do not independently compute whether REPAIR, NO_REPAIR or ESCALATE is the warranted diagnosis-level disposition.

Among the works we reviewed in a search on 30 September 2026, we did not find permission, validity and justification separated into three separately implemented authorities; the separations we found are two-way, or, like Bensalem et al.’s, along other axes.

10. Conclusion

Permission, validity and justification are different questions, and in our data they gave different answers. With tools, gpt-6-sol told apart all twelve twins and all six held-out cases. Across four investigators, admission agreed with the answer key on every model-proposed repair in the registered packs. The ten hosted repairs, each matching the reference fix row for row, were admitted. All 27 of the local model’s proposals were refused, including patches that damaged the data. The checks are effective but not complete. In 6 of the local model’s 9 repaired cases no repair was justified at all; permission offered no resistance, and the validator’s change check refused them. A scripted patch that corrupted every order’s amount while keeping the checked totals passed both checks in every world, so admission does not establish that the rows are correct. The hosted models erred the other way, toward inaction, which no runtime check assesses. The decision to act therefore needs its own evaluation, apart from the checks that admit an intervention.

A clean admission record does not establish entitlement. Under a rule for earned authority, registered before the hosted runs and applied afterwards to the local one, no wrong decision was admitted for any subject or class. That zero count held for both the local model’s all-wrong repairs, which the validator blocked, and the hosted models’ wrong decisions that generated no repair. Correctness scored against the keys distinguished their decisions. Qualification must count both wrong proposals that protection blocks and wrong decisions that produce no proposal for protection to inspect. If authority is to be earned from evidence, that evidence must cover the subject’s decisions,

including those excluded from the admitted set. An authority decision cannot be justified by an observable that fails to identify the competence that authority is intended to reward. The claim is sized to 12 benchmark and 6 held-out cases in one domain, and to a rule that no class could clear.

Future work should test whether the decision to act can itself be challenged before a patch is evaluated, for example by checking the unmodified world before admitting a repair (§7.2). It should also test, separately, cases where the right fix lies outside the actor’s authority and cases where governance has not yet chosen an outcome.

Author contributions

Roles follow the CRediT taxonomy. Malek Alhazmi: conceptualization, methodology, software, validation, investigation, formal analysis, visualization, project administration, writing (original draft; review and editing). Joorie Alsakran: methodology, software, validation, investigation, formal analysis, writing (review and editing). Ibrahim Altowalah: software, validation, writing (review and editing). Nasser Alzaid: software, validation, writing (review and editing).

Use of AI tools

An AI assistant helped draft this manuscript and ran the repository and literature searches recorded in Appendix A. The authors take responsibility for its content.

Artifact availability

The code, incident packages, frozen answer and grounding keys, freeze files, and pack receipts and reports are in the project repository, github.com/AethersDev/adii_team. The code is released under AGPL-3.0-only; the incidents, keys and run archives under CC BY-SA 4.0. The repository state that supports this version of the paper is the tag **paper-v1**.

- **Hosted packs:** commit 736ae43, freeze freeze-2026-09-26 (digest 3d53782df4c5...). Their 66 run folders are committed as `ADII_final_packs.zip`.
- **Local extension:** commit 477cdb5, freeze freeze-2026-09-30 (digest ad2b06956450...). Its receipt, reports and 36 run folders are in `ADII_local_qwen_pack.zip`. The commands that ran it are in `02_src/docs/qwen_local_extension.md`, and `02_src/scripts/layer_counts.py` reproduces Table 5. The system prompt and the text protocol are in `02_src/adii/provider/openai_compatible.py`.
- **Run archive:** `01_data/runs/MANIFEST.json` attests the archived runs by digest, including the scripted square’s and the local extension’s.
- **This paper:** `paper/paper.md` is its source. `python 02_src/scripts/paper_claims.py` verifies both archives against the manifest and recomputes every reported result from them; `02_src/tests/integration/test_the_paper_claims.py` runs it, and shows each check failing on a changed input.

References

- [1] C. J. Clopper and E. S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934. [doi:10.1093/biomet/26.4.404](https://doi.org/10.1093/biomet/26.4.404)

- [2] X. Liu, Y. E. Zhang, V. Kaspova, P. Rabbani, P. S. Zahraei, T. Zhang, A. Ebrahimpour-Boroojeny and V. Chandrasekaran. AgentAbstain: Do LLM agents know when not to act? [arXiv:2607.10059](#), 2026.
- [3] T. Gloaguen, N. Mündler, M. Müller, V. Raychev and M. Vechev. Coding agents don’t know when to act. [arXiv:2605.07769](#), 2026.
- [4] J. Flynt. GroundEval: A deterministic replacement for LLM-as-judge in stateful agent evaluation. [arXiv:2606.22737](#), 2026.
- [5] M. Turk. Counterfactual evaluation reveals hidden capability profiles in clinical LLMs and agents. [arXiv:2605.30590](#), 2026.
- [6] J. Dai and T. Huang. TwinCheck: Evidence-grounded negative-twin verification for stateful tool agents. [arXiv:2609.26911](#), 2026.
- [7] Z. Tao and C. Yin. When tools silently lie: Evaluating and mitigating blind compliance in tool-augmented data agents. [arXiv:2609.37153](#), 2026.
- [8] S. Bensalem, Y. Dong, M. Franzle, X. Huang, J. Kroger, D. Nickovic, A. Nouri, R. Roy and C. Wu. Position: A three-layer probabilistic assume-guarantee architecture is structurally required for safe LLM agent deployment. [arXiv:2605.18672](#), 2026.
- [9] U.S. Food and Drug Administration. Marketing submission recommendations for a predetermined change control plan for artificial intelligence-enabled device software functions. Final guidance, docket FDA-2022-D-2628; issued 4 December 2024, reissued 18 August 2025. [fda.gov](#)
- [10] K. J. K. Feng, D. W. McDonald and A. X. Zhang. Levels of autonomy for AI agents. [arXiv:2506.12469](#), 2025.
- [11] H. Zheng, Q. Dong, R. K. Depena, J. D. Bhatia, F. Xiao and P. Xu. Separating capability from permission: A governance framework for agentic AI autonomy levels. [arXiv:2607.23438](#), 2026.
- [12] R. Taneja and T. Weber. Scan the skill, govern the action: Composing registry verdicts with runtime consequence control. [arXiv:2609.12001v1](#), 2026.
- [13] S. Yu, F. Carroll and B. L. Bentley. The causal impact of tool affordance on safety alignment in LLM agents. [arXiv:2603.20320](#), 2026.
- [14] M. Pacella, G. Papadia and V. Giliberti. Governed agentic process automation: A floor-safety guarantee for compliance-critical LLM routing. *Algorithms*, 19(8):627, 2026. [doi:10.3390/a19080627](#)
- [15] R. Greenblatt, B. Shlegeris, K. Sachan and F. Roger. AI control: Improving safety despite intentional subversion. In *ICML*, 2024. [arXiv:2312.06942](#)
- [16] S. Ray. What can be enforced? A theory of certified runtime safety for tool-using agents. [arXiv:2607.22868](#), 2026.
- [17] M. Schwall, T. Daniel, T. Victor, F. Favarò and H. Hohnhold. Waymo public road safety performance data. [arXiv:2011.00038](#), 2020.
- [18] N. Webb, D. Smith, C. Ludwick, T. Victor, Q. Hommes, F. Favarò, G. Ivanov and T. Daniel. Waymo’s safety methodologies and safety readiness determinations. [arXiv:2011.00054](#), 2020.

Appendix A. Claims ledger

Each row names the evidence behind a claim in the paper: a file in the repository, a run record, or a page opened on 30 September 2026. A claim that something is absent names where we searched. Unrun marks a premise the paper does not state as a finding.

claim	evidence, and where we searched	status
gpt-6-sol right on 12/12 with tools and 4/12 on the alert alone, equal to the always-escalate floor	<code>02_src/docs/evaluation_report.md</code> , generated from the four hosted pack reports	Verified
Hosted packs: 10 repairs proposed, all right and admitted; no wrong fix admitted	<code>evaluation_report.md</code> , Limitations	Verified
All ten hosted repairs were settled by the LLM judge, which ruled each right	each run’s <code>evaluation_report.json</code> names <code>gpt-4.1-mini</code> and its verdict	Verified
The ten hosted repairs rebuild the staged orders and the mart exactly as the reference fix does	each patch rebuilt from frozen inputs and compared row by row with the reference staging, 30 Sep	Verified
Every model run received the same system prompt	the first request’s system message in all 90 model runs, one sha256	Verified
No held-out case ran before its final pack; one retired case from each held-out company ran in rehearsal	run archive; <code>partition.json</code> committed after the rehearsal packs of 23 Sep	Verified
Qwen: repairs in 9/12 cases (27 proposals), all authorized, all rejected, none admitted; right on 2/12 cases (6/36)	<code>01_data/packs/local-qwen3-4b.report.json</code> ; <code>02_src/scripts/layer_counts.py</code>	Verified
6 of Qwen’s 9 repaired cases (18 runs) were no-fix worlds, and every patch there was empty	validation reports in those 18 records	Verified
In all 12 cases, repeats produced identical decisions (patch text included) and tool-call sequences	<code>repeat_disagreements</code> is empty in the pack report; decision and tool-call hashes compared across the 36 records	Verified
The two freezes differ in three files, none in the loop, tools, authorizer, validator or scorer	<code>freeze-2026-09-26.json</code> against <code>freeze-2026-09-30.json</code> : <code>evaluation/grid.py</code> , <code>evaluation/lock.py</code> , <code>runtime/main__.py</code>	Verified
The second freeze brought local runs’ completion bound in line with hosted runs	commit 58bb1cd: before it, local runs sent no completion bound, so the local server’s default of 512 tokens applied; hosted runs used 4,096	Verified
Pack and freeze committed before the first run; branch and tag not on origin at the time	commits 58bb1cd, 477cdb5 at 09:30; first receipt 09:32; <code>git ls-remote</code> , 30 Sep	Verified

claim	evidence, and where we searched	status
The architecture boundaries hold	<code>02_src/tests/architecture/test_boundaries.py</code> , in a full suite of 1,440 passing tests at 477cdb5 on Python 3.12	Verified
Before Qwen, every authorizer or validator refusal in the archives was of a team-scripted proposal	the 13 model REPAIR records marked rejected all read “No independent validator exists yet”; searched <code>01_data/runs</code> (234 records), its 16 September copy (21) and <code>adii_archive_preserved</code> (97); the team’s earlier investigator repository, <code>adii_investigator</code> , holds 79 runs with 23 validated repairs, all passed	Verified
No archived run was both denied and rejected; a unit test covers that cell with a stand-in verdict	the same archives; <code>test_repair_authorization.py</code>	Verified
A no-op repair of a healthy world is rejected by design	<code>test_a_repair_of_a_world_that_was_never_broken_here_changes_nothing</code> , on company A’s NO_REPAIR and ESCALATE cases	Verified
The frozen validator has no check on the unpatched world	<code>02_src/adii/validation/validator.py</code> (<code>validate()</code>) and <code>checks.py</code>	Verified
A patch that changes the world and keeps both data checks is admitted	a scripted probe after the registered runs: every amount replaced by its day’s average, admitted in all 18 live worlds; <code>test_a_patch_that_keeps_every_order_and_every_daily_total_is_admitted</code>	Observed, unregistered
An unpatched-world precondition separates repair worlds from the rest	7 of 12 worlds read through Qwen’s empty patches; the no-op test adds company A’s NO_REPAIR case, asserting only the rejection	Partly observed
Every reported empirical result, recomputed from the published archives	<code>02_src/scripts/paper_claims.py</code> , one identifier per claim; run by <code>test_the_paper_claims.py</code>	Verified
The earned-authority rule was registered before any final run and is never enforced	<code>02_src/docs/evaluation_report.md</code> , “Earned authority”	Verified
Prior work on permission, abstention, grounding and qualification	literature search, 30 Sep (§9); reopened by hand: Feng et al., Zheng et al., Taneja and Weber v1 §8, the FDA guidance’s cover page, Bensalem et al.’s abstract	Verified
Prior work on outcomes behind an interlock	search, 1 Oct (§9); full text read: Yu et al. §IV.B and §VI.B, Schwall et al. §1.2, Webb et al.; abstracts read: Pacella et al., Ray, Greenblatt et al.	Verified
Under the earned-authority rule, no wrong decision was admitted for any subject or class; Qwen’s fix-it class had 0 of 27 right	<code>earned()</code> in <code>02_src/scripts/evaluation_report.py</code> , applied after its runs to <code>local-qwen3-4b.report.json</code> , whose registration named no such count; the hosted table in <code>evaluation_report.md</code>	Verified

claim	evidence, and where we searched	status
A right-disposition REPAIR counts as right only if the validator accepted it, and with the key’s identifiers is scored right without its patch being read	<code>decide_route</code> in <code>02_src/adii/evaluation/scoring.py</code> ; the averaging patch with the key’s identifiers through <code>score_decision</code> on the six repair cases: all accepted, all <code>correct</code> , settled deterministically, 1 Oct	Verified

Appendix B. Case identifiers

Table 8: Frozen incident identifiers of the benchmark cases, by company.

company	tier	REPAIR	NO_REPAIR	ESCALATE
A	explicit	revenue-drop-610a79	revenue-drop-2c077d	revenue-drop-70e933
B	implicit	revenue-drop-982e8e	revenue-drop-d22c8a	revenue-drop-4c1f15
C	explicit	revenue-drop-594c0f	revenue-drop-6fb41f	revenue-drop-70d601
D	implicit	revenue-drop-18b88d	revenue-drop-03df8b	revenue-drop-8e68f4